

Computer Architecture Practical Exercise

0 Cluster Introduction

Kenan Gündogan¹ Philipp Gündisch¹

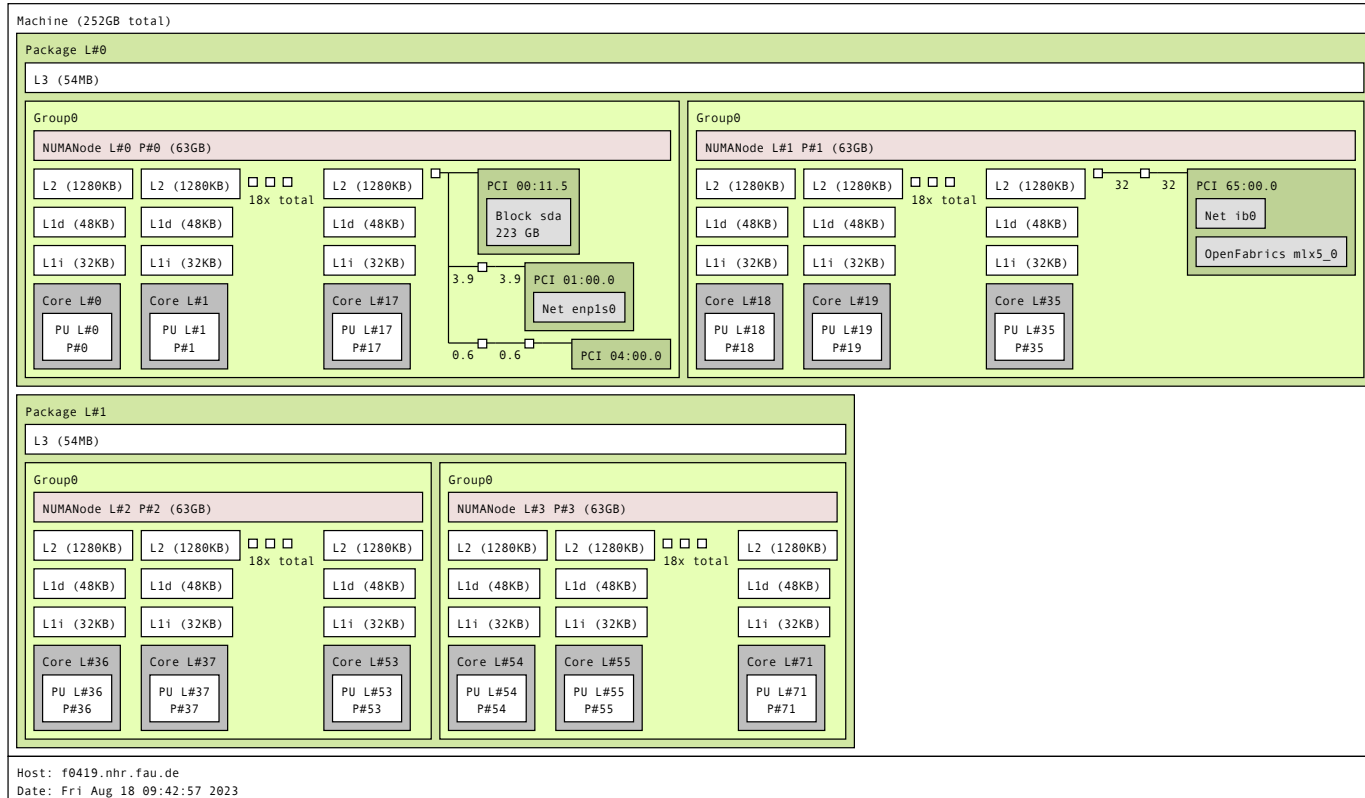
¹Friedrich-Alexander Universität Erlangen-Nürnberg, Chair of Computer Science 3
(Computer Architecture)

October 29, 2024

- The compute nodes of the **NHR** (Erlangen National High Performance Computing Center) are used in this exercise
- The **fritz** cluster consists of over 1000 compute nodes
 - 100 GBit/s HDR100 Infiniband interconnect
 - Dual-socket system with 16 x 16 GB of DDR4-3200 RAM
 - ▶ Ice Lake (*Intel Xeon Platinum 8360Y*)
 - ▶ 36 cores per chip
 - ▶ 2.4 GHz clock frequency
 - More information: see [fritz](#)

- Dual-socket Ice Lake node

- 2x Xeon Platinum 8360Y Chips (CPUs, Packages) with 16 x 16 GB of RAM
- 36 physical cores and 72 logical cores (SMT-Threads) per processor
- CPU clock frequency of 2.4 GHz (nominal) and 3.5 GHz peak (turbo) frequency




- For each student an HPC account (*hpcvXXXv*) was created
- **You should have a mail invitation received**
- There is no password authentication possible
- For authentication to the server an ssh key needs to be generated and uploaded
- For configuration visit the [HPC Portal](#)

Note: The synchronization of the ssh public key to the servers can take up to two hours!

Log in to Cluster

- Access from outside of the university network is possible via dialog-server:

```
ssh hpcvXXXh@fritz.nhr.fau.de
```



- There is no password authentication possible
- From the dialog server / university network the cluster head node is reachable via `ssh`
- To run a program on the cluster `slurm` tools need to be used



salloc

- Slurm is a cluster management and job scheduling system
- Cluster access needs to be requested via Slurm
- A cluster node can be allocated for 1 hour with:

```
salloc
  --partition=singlenode    // fritz subset of nodes
  --time=01:00:00          // allocate for 1 hour
  --nodes=1                 // allocate 1 node
  --ntasks-per-node=1      // for 1 process
  --cpus-per-task=1        // for 1 thread
```

sbatch

The `sbatch` command is able to allocate cluster resources and run a shell script. The resource allocation information is parsed by `sbatch` from the provided shell script (one argument per line).

An exemplary `myjob.sh` file header might look like this:

```
#!/bin/bash -l

#SBATCH --partition=singlenode
#SBATCH --time=00:10:00
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=1
#SBATCH --cpus-per-task=1
#SBATCH --export=NONE

...
srun ./program
```

The shell script can be invoked with:

```
sbatch myjob.sh
```

You can find more information on the [NHR website about slurm](#).

- Use the cluster responsibly!
- Allocate only one node at a time
- Calculate / predict the job runtime and specify `--time`
- Test your implementation and script before submitting a long running job

Shell Functions

You can write functions and invoke them in bash scripts very similar to other programming languages. The following example illustrates how it works.

```
#!/bin/bash -l

function multiply() {
    echo $(( $1*$2 ))
}

echo $(multiply 2 10)
```

When invoked the example produces the following output.

```
$ ./example.sh
20
```

For-Loops

You can also write for-loops in bash scripts. The following example illustrates how for-loops work.

```
#!/bin/bash -l

for name in "Alf" "Joe"
do
    echo $name
done

for (( i = 0; i < 2; i += 1 ))
do
    echo $i
done
```

When invoked the example produces the following output.

```
$ ./example.sh
Alf
Joe
0
1
```

Exercise 0.1: Log in to cluster



- Configure the `ssh` access as described in the [NHR website on SSH](#)
- Use public key authentication
- Log in to the cluster with `ssh fritz`

Exercise 0.2: Running provided code on cluster



- To list available modules on a compute node run:
 - `module avail`
- To make `icx` available on a compute node run:
 - `module load intel`
- To compile the code run `icx -i src/hello.c -o bin/hello`
- Inspect current cluster load
 - [HPC status page](#)
 - To retrieve necessary login data please run `docpw` command on current cluster node
- Update the provided shell script to make use of functions and loops (e.g. run `bin/hello` 5 times)
- Use `sbatch` to run the shell script on the cluster
- Select walltime less than 1h: `Devel` queue (usually shorter waiting times)
- Analyze the produced `sbatch` output

- E 0.1: Set Up
 - Configure ssh access
 - Log in to cluster
- E 0.2: Run code on cluster
 - Copy files via `scp`
 - Start an interactive session via `ssh` (terminate with `exit` or `ctrl+D`)
 - Implement on head node
 - Work on compute node only for benchmarking and testing
 - Compile `src/hello.c` with `icx`
 - Run `sbatch scripts/cluster.sh` on cluster