

Computer Architecture Practical Exercise

4 Vectorization

Kenan Gündogan¹ Philipp Gündisch¹

¹Friedrich-Alexander Universität Erlangen-Nürnberg, Chair of Computer Science 3
(Computer Architecture)

November 11, 2024

SIMD

$$G_{x,y}^{i+1} = \frac{G_{x,y+1}^i + G_{x,y-1}^i + G_{x+1,y}^i + G_{x-1,y}^i}{4}$$

- Jacobi performs the same operation over a huge number of grid cells
- There are no data dependencies between the operations (within a time step simulation)

→ SIMD (Single Instruction Multiple Data) techniques can be applied.

Larger registers paired with specialized instructions (*intrinsics*) can be used to batch process multiple numbers at once.

Vectorization


In contrast to scalar instructions, SIMD instructions execute the same instruction on vectors. The Intel x86 architecture (as a CISC) offers specific instruction set extensions:

- **SSE** (Streaming SIMD Extension)
 - **128 bit** instructions
 - 4 float values per instruction
 - 2 double values per instruction
- **AVX** (Advanced Vector Extension)
 - **256 bit** instructions
 - 8 float values per instruction
 - 4 double values per instruction
- **AVX-512** (Advanced Vector Extension 512 Bit)
 - **512 bit** instructions
 - 16 float values per instruction
 - 8 double values per instruction
- many more ...


Note: these extensions offer further operations for other data types as well.

Vectorization


- SSE** (Streaming SIMD Extension)



- AVX** (Advanced Vector Extension)



- AVX-512** (Advanced Vector Extension)



API

Use Intel's official [intrinsics documentation](#) to identify relevant functions.
To make the intrinsic functions available in C you need to include the header:

```
#include <immintrin.h>
```

Example for the naming convention of intrinsic functions: `_mm256_add_pd`

- Identification of the instruction set
 - SSE with `mm`
 - AVX with `mm256`
 - AVX-512 with `mm512`
- Identification of the operation type with
 - `load`, `store`, `add`, `mul`, `div`, ...
- Identification of scalar vs SIMD with
 - scalar with `s`
 - packed (SIMD) with `p`
- Identification of the data type with
 - float (4 bytes) with `s`
 - double (8 bytes) with `d`

Note: You need to ensure the processor supports the dedicated instruction set extension.

Introduction

Unrolled loops can be manually vectorized as displayed in the following example.

```
#include <immintrin.h>
#include "vec_sum.h"

float vec_sum(const float * restrict array, int32_t length) {
    float      sum          = 0.0f;
    int32_t     remainder    = length % 8;
    __m256      result       = _mm256_setzero_ps();
    __m256      accu;

    #pragma nounroll
    #pragma novector
    for(int32_t i = 0 ; i < length-remainder ; i+=8) {
        accu    = _mm256_loadu_ps(&array[i]);
        result  = _mm256_add_ps(result, accu);
    }

    // Add result and remainder elements into 'sum'
    ...
    return sum;
}
```

Assembler Code

Assembler code resembling the for loop of `vec_sum()` without optimization.

```
..B1.5:
    addss    (%rdi,%rax,4), %xmm0
    incl     %eax
    cmpl     %esi, %eax
    jb       ..B1.5
```

Assembler code resembling the eightfold unrolled loop with SIMD instructions.

```
..B1.5:
    vaddps   (%r8,%rax,4), %ymm0, %ymm0
    addl     $8, %eax
    cmpl     %edi, %eax
    jb       ..B1.5
```

Note: the assembler codes were produced with:

```
icx -I ./include/ -S -O3 vec_sum.c
```

Task 4.1: Jacobi SSE & AVX



Implementation & Visualization

- Keep suppressing implicit unrolling and vectorization
- Keep the naive implementation without unrolling
- Adapt the twofold unrolled loop to use SSE instructions
- Adapt the fourfold unrolled loop to use AVX instructions
- Adapt the eightfold unrolled loop to use AVX512 instructions
- Benchmark all three implementations for 1KiB - 128MiB (1 second runtime)
- Draw a line chart with the tool of your choice
- Choose the memory consumption as the X-axis
- Choose the performance metrics for the Y-axis (MUp/s)
- **NOTE: USE UNALIGNED LOADS**

Optional Task 4.2: VecSum



NOT MANDATORY

Beat the compiler by exceeding the performance of e00! (for at last one measurement)

- Keep suppressing implicit unrolling and vectorization
- Adapt the eightfold unrolled loop to use AVX instructions for summation
- Keep the other implementations from previous exercise
- Implement a twofold, threefold and fourfold unrolling for the AVX instructions by replacing `#pragma nounroll` with `#pragma unroll (N)`
- Benchmark all three implementations for 1KiB - 32MiB (1 second runtime)
- Draw a line chart with the tool of your choice
- Choose the memory consumption as the X-axis
- Choose the performance metrics for the Y-axis (AdditionsPerSecond)
- **NOTE: USE UNALIGNED LOADS**

- E 4.1: Jacobi SSE & AVX
 - Update twofold unrolling with SSE instructions
 - Update fourfold unrolling with AVX instructions
 - Compare results in a linechart
- Optional 4.2: VecSum AVX
 - Update eightfold unrolling with AVX instructions
 - Unroll the eightfold unrolled with AVX instructions further
 - Compare results in a linechart

Appendix: CPU Flags



Identify Supported Extensions

There are multiple ways to identify the instruction set extensions of your processor.

- `cat /proc/cpuinfo`
- `lscpu`
- Internet search for the processor model

Performance Optimization

During the timeline of this class new bullet points will be added. Recently added entries are bold.

- Compiling
 - Choice of the compiler (`icx`)
 - Compiler flag to optimize aggressively (e.g. `-O3`)
 - Compiler flag to adapt for specific hardware (e.g. `-xHost`)
- Programming Techniques (if applicable)
 - Use `#define` and `const` instead of variables
 - Data type aware programming
 - Use aligned memory (e.g. with `_mm_malloc()` or `posix_memalign()`)
 - Consecutive address iteration
 - **Variable declarations outside of loops**
 - **Reduce function calls**
 - **Use intrinsics (to utilize SIMD)**
- Measurement
 - Reasonable benchmark time
 - Reasonable benchmark workload
 - Reduce interference factors to a minimum